

Assignment -3

Assignment Date	21 October 2022
Student Name	MEENAKSHI SUNDARAM P
Student Roll Number	210519205030
Maximum Marks	2 Marks

Data Visualization and Pre-processing

Building a Regression Model

1. Perform Below Visualizations.

Univariate Analysis

1. Summary Statistics

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
```

```
In [5]: file_data = pd.read_csv(r'C:\Users\Guru\Desktop\abalone\abalone.csv')
file_data
```

Out[5]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4177 rows x 9 columns

Add a Age column in a dataset

```
In [6]: file_data['Age']=''
file_data.head()
```

Out[6]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7	

```
In [7]: file_data['Age']=file_data['Rings']+1.5
file_data.head()
```

Out[7]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7	8.5

Drop the Rings Column

```
In [8]: file_data = file_data.drop(columns=['Rings'],axis=1)
file_data
```

Out[8]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	8.5
...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	12.5
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	11.5
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	10.5
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	11.5
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	13.5

4177 rows × 9 columns

```
In [9]: file_data['Height'].mean()
```

Out[9]: 0.1395163993296614

```
In [10]: file_data['Height'].median()
```

Out[10]: 0.14

```
In [11]: file_data['Height'].std()
```

```
Out[11]: 0.04182705660725703
```

2. Frequency Table


```
In [12]: file_data['Sex'].value_counts()
```

```
Out[12]: M      1528  
         I      1342  
         F      1307  
         Name: Sex, dtype: int64
```

3. Create Charts


```
In [13]: file_data.boxplot(column=['Length'], grid=False)
```

```
Out[13]: <AxesSubplot:>
```




```
In [14]: file_data.hist(column='Length', grid=False, edgecolor='black')
```

```
Out[14]: array([[<AxesSubplot:title={'center':'Length'}>]], dtype=object)
```



```
In [15]: sns.kdeplot(file_data['Length'])
```

```
Out[15]: <AxesSubplot:xlabel='Length', ylabel='Density'>
```




Bi - Variate Analysis

1. Barplot

```
In [17]: data = sns.barplot(x = file_data["Sex"], y = file_data["Age"])
data
```

```
Out[17]: <AxesSubplot:xlabel='Sex', ylabel='Age'>
```



2. Correlation Coefficients

```
In [18]: file_data.corr()
```


```
Out[18]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
Length	1.000000	0.986812	0.827554	0.925261	0.897914	0.903018	0.897706	0.556720
Diameter	0.986812	1.000000	0.833684	0.925452	0.893162	0.899724	0.905330	0.574660
Height	0.827554	0.833684	1.000000	0.819221	0.774972	0.798319	0.817338	0.557467
Whole weight	0.925261	0.925452	0.819221	1.000000	0.969405	0.966375	0.955355	0.540390
Shucked weight	0.897914	0.893162	0.774972	0.969405	1.000000	0.931961	0.882617	0.420884
Viscera weight	0.903018	0.899724	0.798319	0.966375	0.931961	1.000000	0.907656	0.503819
Shell weight	0.897706	0.905330	0.817338	0.955355	0.882617	0.907656	1.000000	0.627574
Age	0.556720	0.574660	0.557467	0.540390	0.420884	0.503819	0.627574	1.000000

3.Linear Plot

```
In [19]: data = sns.lineplot(x = file_data["Length"], y = file_data["Age"])  
data
```


```
Out[19]: <AxesSubplot:xlabel='Length', ylabel='Age'>
```



4. Scatter Plot

```
In [20]: data = sns.scatterplot(x = file_data['Age'], y = file_data['Diameter'], color="green")  
data
```

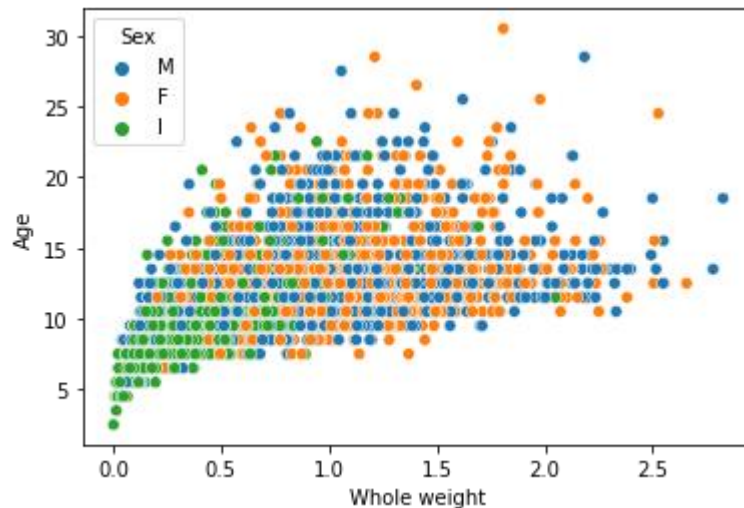
```
Out[20]: <AxesSubplot:xlabel='Age', ylabel='Diameter'>
```



Multi - Variate Analysis

```
In [21]: x = sns.scatterplot(x=file_data['Whole weight'],y=file_data['Age'],hue=file_data["Sex"])
x
```

```
Out[21]: <AxesSubplot:xlabel='Whole weight', ylabel='Age'>
```



4. Perform descriptive statistics on the dataset

```
In [22]: file_data.shape
```

```
Out[22]: (4177, 9)
```

```
In [23]: file_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight   4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Age              4177 non-null   float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
```


In [24]: `file_data.describe()`

Out[24]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell we
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005

In [25]: `file_data.head()`

Out[25]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

In [26]: `file_data.tail()`

Out[26]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	12.5
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	11.5
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	10.5
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	11.5
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	13.5

In [27]: `file_data.mean(numeric_only=True)`

Out[27]:

Length	0.523992
Diameter	0.407881
Height	0.139516
Whole weight	0.828742
Shucked weight	0.359367
Viscera weight	0.180594
Shell weight	0.238831
Age	11.433684
dtype:	float64

In [28]: `file_data.median(numeric_only=True)`

Out[28]:

Length	0.5450
Diameter	0.4250
Height	0.1400
Whole weight	0.7995
Shucked weight	0.3360
Viscera weight	0.1710
Shell weight	0.2340
Age	10.5000
dtype:	float64

In [29]: `file_data.mode()`

Out[29]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.550	0.45	0.15	0.2225	0.175	0.1715	0.275	10.5
1	NaN	0.625	NaN	NaN	NaN	NaN	NaN	NaN	NaN

In [30]: `file_data.var(numeric_only=True)`

Out[30]:

Length	0.014422
Diameter	0.009849
Height	0.001750
Whole weight	0.240481
Shucked weight	0.049268
Viscera weight	0.012015
Shell weight	0.019377
Age	10.395266
dtype:	float64

```
In [31]: file_data.std(numeric_only=True)
```

```
Out[31]: Length          0.120093
Diameter          0.099240
Height           0.041827
Whole weight      0.490389
Shucked weight    0.221963
Viscera weight    0.109614
Shell weight      0.139203
Age               3.224169
dtype: float64
```

```
In [32]: file_data.skew(numeric_only=True)
```

```
Out[32]: Length          -0.639873
Diameter          -0.609198
Height            3.128817
Whole weight      0.530959
Shucked weight    0.719098
Viscera weight    0.591852
Shell weight      0.620927
Age               1.114102
dtype: float64
```

```
In [33]: file_data.kurt(numeric_only=True)
```


```
Out[33]: Length          0.064621
Diameter          -0.045476
Height           76.025509
Whole weight      -0.023644
Shucked weight    0.595124
Viscera weight    0.084012
Shell weight      0.531926
Age               2.330687
dtype: float64
```

```
In [34]: quantile = file_data['Whole weight'].quantile(q=[0.75, 0.25])
quantile
```

```
Out[34]: 0.75    1.1530
0.25    0.4415
Name: Whole weight, dtype: float64
```

```
In [35]: x = file_data.Diameter
sns.boxplot(x=x)
```

```
Out[35]: <AxesSubplot:xlabel='Diameter'>
```



5. Handle the Missing values

```
In [36]: print(file_data.isnull())
```

	Sex	Length	Diameter	Height	Whole	weight	Shucked	weight \
0	False	False	False	False		False		False
1	False	False	False	False		False		False
2	False	False	False	False		False		False
3	False	False	False	False		False		False
4	False	False	False	False		False		False
...
4172	False	False	False	False		False		False
4173	False	False	False	False		False		False
4174	False	False	False	False		False		False
4175	False	False	False	False		False		False
4176	False	False	False	False		False		False

	Viscera	weight	Shell	weight	Age
0		False		False	False
1		False		False	False
2		False		False	False
3		False		False	False
4		False		False	False
...	
4172		False		False	False
4173		False		False	False
4174		False		False	False
4175		False		False	False
4176		False		False	False

```
[4177 rows x 9 columns]
```

```
In [37]: print(file_data.isnull().sum())
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Age          0
dtype: int64
```


```
In [38]: file_data.isna().any()
```

```
Out[38]: Sex          False
Length       False
Diameter     False
Height       False
Whole weight False
Shucked weight False
Viscera weight False
Shell weight False
Age          False
dtype: bool
```

6. Find the outliers and replace the outliers


```
In [40]: x = sns.boxplot(x=file_data["Age"])
x
```

```
Out[40]: <AxesSubplot:xlabel='Age'>
```



```
In [41]: x = file_data.Age  
sns.boxplot(x=x)
```


```
Out[41]: <AxesSubplot:xlabel='Age'>
```



```
In [42]: x = np.where(file_data['Age']>57,39, file_data['Age'])
```

```
In [43]: sns.boxplot(x=x)
```

```
Out[43]: <AxesSubplot:>
```



7. Check for Categorical columns and perform encoding.

```
In [44]: import warnings
warnings.filterwarnings('ignore')
x = pd.Categorical(file_data["Whole weight"])
x
```

```
Out[44]: [0.5140, 0.2255, 0.6770, 0.5160, 0.2050, ..., 0.8870, 0.9660, 1.1760, 1.094
5, 1.9485]
Length: 4177
Categories (2429, float64): [0.0020, 0.0080, 0.0105, 0.0130, ..., 2.5550,
2.6570, 2.7795, 2.8255]
```

```
In [45]: pd.get_dummies(file_data["Height"]).head(10)
```

```
Out[45]:
```

	0.000	0.010	0.015	0.020	0.025	0.030	0.035	0.040	0.045	0.050	...	0.210	0.215	0.220	0
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
5	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
6	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
7	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
8	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
9	0	0	0	0	0	0	0	0	0	0	...	0	0	0	

10 rows × 51 columns

```
In [46]: pd.get_dummies(file_data).head(10)
```

```
Out[46]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5	0	0	1
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5	0	0	1
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5	0	0	1
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5	0	1	0
5	0.425	0.300	0.095	0.3515	0.1410	0.0775	0.120	9.5	0	1	0
6	0.530	0.415	0.150	0.7775	0.2370	0.1415	0.330	21.5	1	0	0
7	0.545	0.425	0.125	0.7680	0.2940	0.1495	0.260	17.5	1	0	0
8	0.475	0.370	0.125	0.5095	0.2165	0.1125	0.165	10.5	0	0	1
9	0.550	0.440	0.150	0.8945	0.3145	0.1510	0.320	20.5	1	0	0

8. Split the data into dependent and independent variables.

```
In [48]: X = file_data.iloc[:, :-1].values
print(X)
```

```
[[ 'M' 0.455 0.365 ... 0.2245 0.101 0.15]
 [ 'M' 0.35 0.265 ... 0.0995 0.0485 0.07]
 [ 'F' 0.53 0.42 ... 0.2565 0.1415 0.21]
 ...
 [ 'M' 0.6 0.475 ... 0.5255 0.2875 0.308]
 [ 'F' 0.625 0.485 ... 0.531 0.261 0.296]
 [ 'M' 0.71 0.555 ... 0.9455 0.3765 0.495]]
```

```
In [49]: Y = file_data.iloc[:, -1].values
print(Y)
```

```
[16.5  8.5 10.5 ... 10.5 11.5 13.5]
```

9. Scale the independent variables

```
In [50]: from sklearn.preprocessing import scale
```



```
In [54]: x = scale(file_data["Viscera weight"])
x
```

```
Out[54]: array([-0.72621157, -1.20522124, -0.35668983, ...,  0.97541324,
                0.73362741,  1.78744868])
```

10. Split the data into training and testing

```
In [55]: from sklearn.model_selection import train_test_split
```

```
In [56]: x = file_data.iloc[:, 1:7]
x
```

Out[56]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	0.330	0.255	0.080	0.2050	0.0895	0.0395
...
4172	0.565	0.450	0.165	0.8870	0.3700	0.2390
4173	0.590	0.440	0.135	0.9660	0.4390	0.2145
4174	0.600	0.475	0.205	1.1760	0.5255	0.2875
4175	0.625	0.485	0.150	1.0945	0.5310	0.2610
4176	0.710	0.555	0.195	1.9485	0.9455	0.3765

4177 rows x 6 columns

```
In [57]: y = file_data.iloc[:, -1]
y
```

```
Out[57]: 0      16.5
1       8.5
2      10.5
3      11.5
4       8.5
...
4172    12.5
4173    11.5
4174    10.5
4175    11.5
4176    13.5
Name: Age, Length: 4177, dtype: float64
```

```
In [58]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state =42)
```

```
In [59]: x_train
```

```
Out[59]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
3823	0.615	0.455	0.135	1.0590	0.4735	0.2630
3956	0.515	0.395	0.140	0.6860	0.2810	0.1255
3623	0.660	0.530	0.175	1.5830	0.7395	0.3505
0	0.455	0.365	0.095	0.5140	0.2245	0.1010
2183	0.495	0.400	0.155	0.8085	0.2345	0.1155
...
3444	0.490	0.400	0.115	0.5690	0.2560	0.1325
466	0.670	0.550	0.190	1.3905	0.5425	0.3035
3092	0.510	0.395	0.125	0.5805	0.2440	0.1335
3772	0.575	0.465	0.120	1.0535	0.5160	0.2185
860	0.595	0.475	0.160	1.1405	0.5470	0.2310

3132 rows x 6 columns

```
In [60]: y_train
```

```
Out[60]: 3823    10.5
3956    13.5
3623    11.5
0       16.5
2183     7.5
...
3444    10.5
466     13.5
3092    12.5
3772    10.5
860     7.5
Name: Age, Length: 3132, dtype: float64
```

```
In [61]: print(x_train.shape, x_test.shape)
```

```
(3132, 6) (1045, 6)
```

11. Build the Model

```
In [62]: from sklearn.linear_model import LinearRegression
```

```
In [63]: model=LinearRegression()
```

```
In [64]: model.fit(x_train,y_train)
```

```
Out[64]: LinearRegression()
```

12. Train the Model

```
In [65]: Y_predict_train = model.predict(x_train)
Y_predict_train
```

```
Out[65]: array([11.25888828, 11.95379472, 12.33692259, ..., 11.12903068,
                10.71152746, 11.59516371])
```

13. Test the Model

```
In [66]: y_predict = model.predict(x_test)
y_predict
```

```
Out[66]: array([13.0478407 , 11.43166184, 15.59825921, ..., 13.69440346,
                11.79279231, 10.83037939])
```

14. Measure the performance using Metrics

```
In [67]: from sklearn.metrics import mean_squared_error
import math
print(mean_squared_error(y_test, y_predict))
print(math.sqrt(mean_squared_error(y_test, y_predict)))
```

```
4.862459933051859
2.205098622069285
```

```
In [ ]:
```